

Architecture Overview – Networking – ML model development/update

1. User enters Kubeflow FQDN in the web browser and is forwarded to the Ingress controller
2. Ingress controller translates the FQDN to the appropriate service and forwards it
3. Kubeflow GUI comes up in the Users web browser and user can spawn a Jupyter Notebook in the Kubeflow to develop a new ML inference model

Service name	Service type	Note
Vehicle API	NodePort	externalTrafficPolicy: Local Externally Accessible through Anycast/NodeIP:nodePort
Kserve Inference	ClusterIP	internalTrafficPolicy: Local
Kepler	ClusterIP	
Prometheus	ClusterIP	
Kserve	ClusterIP	
Kubeflow	ClusterIP	GUI accessible through Ingress
Minio Storage	ClusterIP	
Grafana	ClusterIP	GUI accessible through Ingress

